


Lecture 2. Amortized Analysis

CpSc 8400: Algorithms and Data Structures
Brian C. Dean



School of Computing
Clemson University
Spring, 2021

1


Warm-Up: Fundamental Data Structures

Arrays



- Retrieve or modify any element in $O(1)$ time.
- Insert or delete in middle of list: $O(N)$ time. ☹️
- Insert or delete from ends: $O(1)$ time
 - Be careful not to run over end of allocated memory

Linked Lists



(sometimes doubly linked, or ending with a sentinel instead of NULL)


- Seek to any position in list: $O(N)$ time. ☹️
- Then insert or delete element: $O(1)$ time.
- Insert or delete from ends: $O(1)$ time.

2

2

Fundamental Data Structures

Queues



- First-In, First-Out (FIFO).
- $O(1)$ enqueue & dequeue, implemented using arrays or linked lists (often implemented using circular arrays).

Stacks



- Last-In, First-Out (LIFO).
- $O(1)$ push & pop, implemented using arrays or linked lists.

3

3

An Important Distinction...

Specification of a data structure in terms of the operations it needs to support.

(sometimes called an *abstract data type*)

A concrete approach for **implementation** of the data structure that fulfills these requirements.

4

4



Example: Queues

Abstract data type:
queue

Must support these operations:

- *Insert(k)* a new key *k* into the structure.
- *Remove* the least-recently-inserted key from the structure. (so FIFO behavior)

Choices for concrete implementation:



5

5

Enforcing Abstraction in Code

Abstract data type:
queue

Concrete implementation:
queue.cpp

queue.h:

```
class Queue {
private:
    int *A;
    int front, back, N;

public:
    Queue();
    ~Queue();
    void insert(int key);
    int remove(void);
};
```

←

→

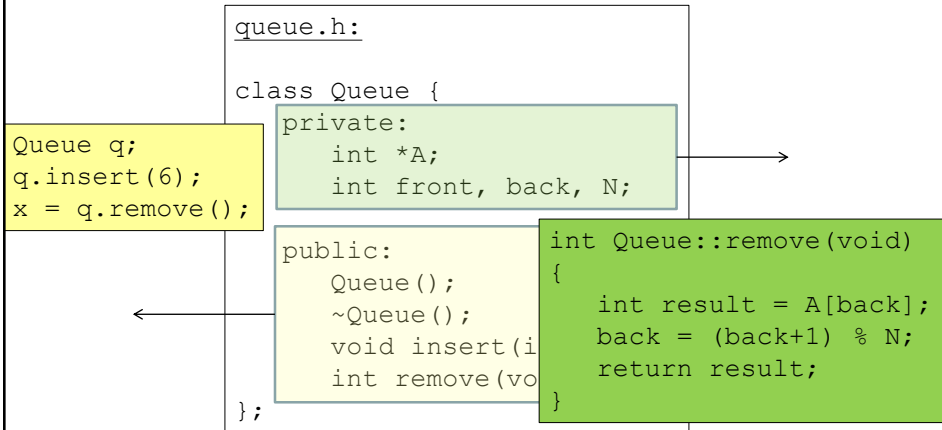
6

6

Enforcing Abstraction in Code

Abstract data type:
queue

Concrete implementation:
queue.cpp



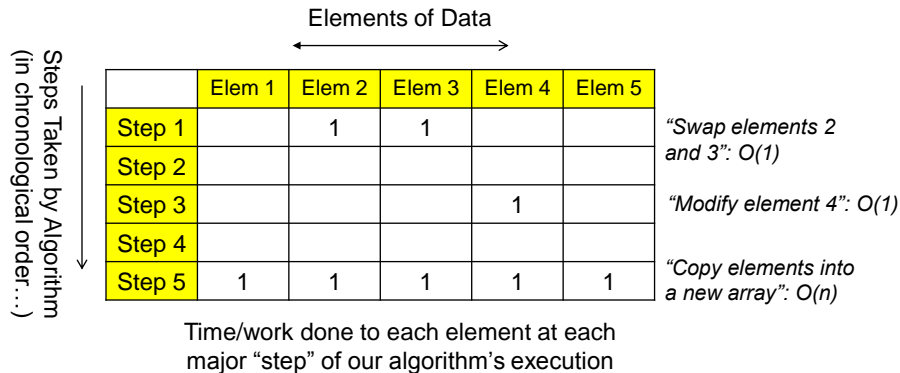
7

Today's Useful Analysis Technique: Adding Things in Smart Ways

- Lots of analyses get easier when you add things together after re-grouping them in smart ways.

8

Example: Think about an Algorithm from the Perspective of a Data Element...





(that is, sum the columns first, not the rows...)

9

Example: Merge Sort


- It takes $\Theta(n)$ time to merge two sorted lists of combined length n .
- How much time does it take to do a full merge sort then?



10

Example: Merge Sort

- It takes $\Theta(n)$ time to merge two sorted lists of combined length n .
- Think of this as $O(1)$ time per element taking part in the merge.
- Now look at a particular array element e . The total “work” we spend on e is equal to the number of merges in which e takes part: $O(\log n)$.
 - Why $O(\log n)$? Each merge doubles the size of the sorted subarray containing e .
- So $O(n \log n)$ total work.



11

11

Useful Analysis Technique: Think about an Algorithm from the Perspective of a Data Element...

- Figure out how much work / running time is spent on a single generic element of data during the course of the algorithm.
- Add this up to get the total running time.
(compared to adding up the time spent on each “operation”, summed over each operation in chronological order)

12

12

Example: Enumerating Subsets

counter = 0

For all subsets $S \subseteq \{1, 2, 3, \dots, n\}$

Increment counter

What is the value of counter at the end of execution?

13

13

Example: Enumerating Subsets

counter = 0

For all subsets $S \subseteq \{1, 2, 3, \dots, n\}$

For all subsets $T \subseteq S$


Increment counter

What is the value of counter at the end of execution?

14

14


Example: Domination Radius



- Given the heights of N individuals standing in a line.
- **Goal:** find the domination radius of each individual.

15


Example: Domination Radius



- Given the heights of N individuals standing in a line.
- **Goal:** find the domination radius of each individual.
- Simple algorithm: from each element, scan left until blocked, then scan right until blocked.

16


Example: Domination Radius



- Given the heights of N individuals standing in a line.
- **Goal:** find the domination radius of each individual.
- Simple algorithm: from each element, scan left until blocked, then scan right until blocked. ($O(N^2)$ worst-case)

17


Example: Domination Radius



- Given the heights of N individuals standing in a line.
- **Goal:** find the domination radius of each individual.
- Simple algorithm: from each element, scan left until blocked, then scan right until blocked. ($O(N^2)$ worst-case)
- Refinement: from each element, scan left and right simultaneously until blocked.

18

Example: Domination Radius



- Given the heights of N individuals standing in a line.
- **Goal:** find the domination radius of each individual.
- Simple algorithm: from each element, scan left until blocked, then scan right until blocked. ($O(N^2)$ worst-case)
- Refinement: from each element, scan left and right simultaneously until blocked.
 - Add up running time by grouping together “high work” elements separately from “low work” elements...

19

Recall: Think about an Algorithm from the Perspective of a Data Element...

		Elements of Data					
		Elem 1	Elem 2	Elem 3	Elem 4	Elem 5	
Steps Taken by Algorithm (in chronological order...)	Step 1		1	1			“Swap elements 2 and 3”
	Step 2						
	Step 3				1		“Modify element 4”
	Step 4						
	Step 5	1	1	1	1	1	“Copy elements into a new array”

Time/work done to each element at each major “step” of our algorithm’s execution

20

20

Re-Sizing Memory Blocks

- Since memory blocks often cannot expand after allocation, what do we do when a memory block fills up?
- For example, suppose we allocate 100 words of memory space for a stack (implemented as an array), but then realize we have more than 100 elements to push onto the stack!

(yes, use of a linked list would have solved this problem, but suppose we really want to use arrays instead...)

21

21

Memory Allocation : Successive Doubling

- A common technique for block expansion: whenever our current block fills up, allocate a new block of twice its size and transfer the contents to the new block.
- Unfortunately, now some of our push operations will be quite slow!
 - Most push operations take only $O(1)$ time.
 - However, a push operation resulting in an expansion (and a copy of the n elements currently in the stack) will take $\Theta(n)$ time.

22

22

How to Describe the Running Time of Push...?

- Push has a somewhat non-uniform running time profile:
 - $O(1)$ almost always
 - Except $\Theta(N)$ every now and then.
- But just saying the running time is “ $\Theta(N)$ in the worst case” doesn’t tell the whole story...
 - Doesn’t do the structure justice.
 - People might be scared to use it for large input sizes...

23

23

How Expensive is Your Car to Maintain...?

Jan: \$10

Feb: \$10

...

Nov: \$10

Dec: \$130 = \$10 + yearly \$120 tune-up

- Same problem: saying it’s “\$130/month in the worst case” doesn’t tell the complete story...

24

24

Amortization of Costs

Actual Cost

Jan: \$10

Feb: \$10

...

Nov: \$10

Dec: \$130

Amortized Cost

Jan: \$20

Feb: \$20

...

Nov: \$20

Dec: \$20

- So our cost is “\$20/month, amortized”.
- This is a simpler, more accurate description of our cost structure.
- Compare with actually paying \$20/month...

25

25

Back to Push...

- How much does each push actually cost?

Op#:	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	...
Insert:	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	...
Copy:		1	2		4				8								16		...
Total:	1	2	3	1	5	1	1	1	9	1	1	1	1	1	1	1	17	1	...
Cumulative:	1	3	6	7	12	13	14	15	24	25	26	27	28	29	30	31	48	49	...

26

26

Back to Push...

- How much does each push actually cost?

Op#:	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	...
Insert:	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	...
Copy:		1	2		4				8								16		...
Total:	1	2	3	1	5	1	1	1	9	1	1	1	1	1	1	1	17	1	...
Cumulative:	1	3	6	7	12	13	14	15	24	25	26	27	28	29	30	31	48	49	...

- What about if we charge ourselves 3 units of work per operation instead...?

Total:	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	...
Cumulative:	3	6	9	12	15	18	21	24	27	30	33	36	39	42	45	48	51	54	...

27

27

Back to Push...

- How much does each push actually cost?

Op#:	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	...
Insert:	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	...
Copy:		1	2		4				8								16		...
Total:	1	2	3	1	5	1	1	1	9	1	1	1	1	1	1	1	17	1	...
Cumulative:	1	3	6	7	12	13	14	15	24	25	26	27	28	29	30	31	48	49	...

- What about if we charge ourselves 3 units of work per operation instead...?

Total:	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	...
Cumulative:	3	6	9	12	15	18	21	24	27	30	33	36	39	42	45	48	51	54	...

“True” cumulative cost after any sequence of k operations is upper bounded by “fictitious” cumulative cost of $3k$...

28

28

Back to Push...

- How much does each push actually cost?

Op#:	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	...
Insert:	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	...
Copy:		1	2		4				8								16		...
Total:	1	2	3	1	5	1	1	1	9	1	1	1	1	1	1	1	17	1	...
Cumulative:	1	3	6	7	12	13	14	15	24	25	26	27	28	29	30	31	48	49	...

- What about if we charge ourselves 3 units of work per operation instead...?

Total:	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	...
Cumulative:	3	6	9	12	15	18	21	24	27	30	33	36	39	42	45	48	51	54	...

Total insert cost = k

Total copy cost $\leq 2k$

29

29

Back to Push...

- How much does each push actually cost?

Op#:	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	...
Insert:	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	...
Copy:		1	2		4				8	+2	+2	+2	+2	+2	+2	+2	16		...
Total:	1	2	3	1	5	1	1	1	+2	1	1	1	1	1	1	1	17	1	...
Cumulative:	1	3	6	7	12	13	14	15	24	25	26	27	28	29	30	31	48	49	...

- What about if we charge ourselves 3 units of work per operation instead...?

Total:	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	...
Cumulative:	3	6	9	12	15	18	21	24	27	30	33	36	39	42	45	48	51	54	...

Total insert cost = k

Total copy cost $\leq 2k$

30

30

Back to Push...

- How much does each push actually cost?

Op#:	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	...
Insert:	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	...
Copy:		1	2		4	+2	+2	+2	8	+2	+2	+2	+2	+2	+2	+2	16		...
Total:	1	2	3	1	+2	1	1	1	+2	1	1	1	1	1	1	1	17	1	...
Cumulative:	1	3	6	7	12	13	14	15	24	25	26	27	28	29	30	31	48	49	...

- What about if we charge ourselves 3 units of work per operation instead...?

Total:	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	...
Cumulative:	3	6	9	12	15	18	21	24	27	30	33	36	39	42	45	48	51	54	...

Total insert cost = k

Total copy cost $\leq 2k$

31

31

Back to Push...

- How much does each push actually cost?

Op#:	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	...
Insert:	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	...
Copy:		1	2		4				8								16		...
Total:	1	2	3	1	5	1	1	1	9	1	1	1	1	1	1	1	17	1	...
Cumulative:	1	3	6	7	12	13	14	15	24	25	26	27	28	29	30	31	48	49	...

- What about if we charge ourselves 3 units of work per operation instead...?

Total:	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	...
Cumulative:	3	6	9	12	15	18	21	24	27	30	33	36	39	42	45	48	51	54	...

- So how different is our version of push from a version that takes 3 units in the worst case?

32

32

Amortized Analysis

- Any sequence of k pushes takes $O(k)$ worst-case time, so we say that push takes $O(1)$ **amortized** time.
- “On average”, over the entire sequence, each individual push therefore takes $O(1)$ time.
- In general, an operation runs in $O(f(n))$ amortized time if any sequence of k such operations runs in $O(k f(n))$ time.

33

33

Amortized Analysis : Motivation

- Amortized analysis is an ideal way to characterize the worst-case running time of operations with highly non-uniform performance.
- It is still **worst-case analysis, just averaged over an arbitrary sequence of operations.**
- It gives us a much clearer picture of the true performance of a data structure that more faithfully describes the true performance.
 - E.g., “ $\Theta(N)$ worst case vs. $O(1)$ amortized”.

34

34

Amortized Analysis : Motivation

- Suppose we have 2 implementations of a data structure to choose from:
 - A: $O(\log n)$ worst-case time / operation.
 - B: $O(\log n)$ amortized time / operation.
- There is **no difference** if we use either A or B as part of a larger algorithm. For example, if our algorithm makes n calls to the data structure, the running time is $O(n \log n)$ in either case.
- The choice between A and B only matters in a “real-time” setting when the response time of an individual operation is important.

35