

# Lecture 5. Priority Queues and Randomized Algorithm Analysis

CpSc 8400: Algorithms and Data Structures  
Brian C. Dean



School of Computing  
Clemson University  
Spring, 2021

1

## Priority Queues

- In a simple FIFO queue, elements exit in the same order as they enter.
- In a priority queue, the element with highest priority (usually defined as having *lowest* key) is always the first to exit.
- Many uses:
  - **Scheduling:** Manage a set of tasks, where you always perform the highest-priority task next.
  - **Sorting:** Insert  $n$  elements into a priority queue and they will emerge in sorted order.
  - **Complex Algorithms:** For example, Dijkstra's shortest path algorithm is built on top of a priority queue.

2

2

## Priority Queues

- All priority queues support:

*Insert*( $e, k$ ) : Insert a new element  $e$  with key  $k$ .

*Remove-Min* : Remove and return the element with minimum key.

- In practice (mostly due to Dijkstra's algorithm), many support:

*Decrease-Key*( $e, \Delta k$ ) : Given a pointer to element  $e$  within the heap, reduce  $e$ 's key by  $\Delta k$ .

- Some priority queues also support:

*Increase-key*( $e, \Delta k$ ) : Increase  $e$ 's key by  $\Delta k$ .

*Delete*( $e$ ) : Remove  $e$  from the structure.

*Find-min* : Return a pointer to the element with minimum key.

3

3

## Redundancies Among Operations

- Given *insert* and *delete*, we can implement *increase-key* and *decrease-key*.
- Given *decrease-key* and *remove-min*, we can implement *delete*.
- Given *find-min* and *delete*, we can implement *remove-min*.
- Given *insert* and *remove-min*, we can implement *find-min*.

4

4

## Priority Queue Implementations

- There are *many* simple ways to implement the abstract notion of a priority queue as a concrete data structure:

	<i>insert</i>	<i>remove-min</i>
Unsorted array or linked list	$O(1)$	$O(n)$
Sorted array or linked list	$O(n)$	$O(1)$
Binary heap	$O(\log n)$	$O(\log n)$
Balanced binary search tree	$O(\log n)$	$O(\log n)$
Skew heap	$O(\log n)$ am.	$O(\log n)$ am.
Randomized mergeable binary heap	$O(\log n)$ whp.	$O(\log n)$ whp.
Binomial heap	$O(1)$ am.	$O(\log n)$
Fibonacci heap	$O(1)$ am.	$O(\log n)$ am.

5

5


## Warm-Up: Incremental Priority Queues

- Fundamental operations of a (max-) priority queue:
  - Insert* : insert new element
  - Remove-max* : remove element with maximum key
- We'll study general priority queues in a moment, but for now, consider the special case of an *incremental* priority queue:
  - Keys stored in the structure are nonnegative integers, initially zero.
  - We additionally support an *increment-priority(e)* operation that takes a pointer to an element and increases its key by 1.

6

6

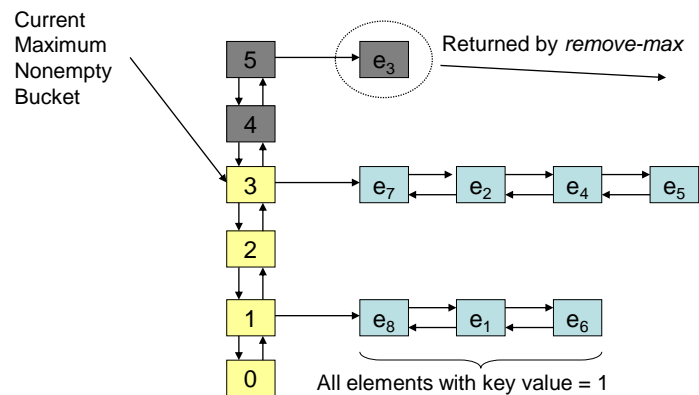
## Implementing an Incremental Priority Queue



7

7

## The Remove-Max Operation



8

8

## Analysis of Incremental Priority Queue

- Let  $M$  denote the amount by which the “current maximum bucket” pointer moves.

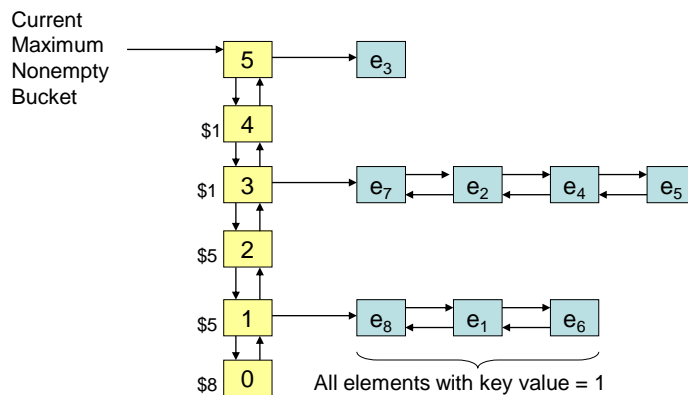
	Worst-Case Running Time	Amortized Running Time
<i>insert</i>	1	1
<i>increment-priority</i>	1	2
<i>remove-max</i>	$1+M$ (not bounded!)	1

↑  
All operations have  $O(1)$  amortized running times!

9

9

## Plenty of Credit to Spare...





10

10

## The Binary Heap

- An almost-complete binary tree (all levels full except the last, which is filled from the left side up to some point).
- Satisfies the **heap property**: for every element  $e$ ,  $\text{key}(\text{parent}(e)) \leq \text{key}(e)$ .
  - Minimum element always resides at root.
- Physically stored in an array  $A[0 \dots n-1]$ .
- Easy to move around the array in a treelike fashion:
  - $\text{Parent}(i) = \text{floor}((i-1)/2)$ .
  - $\text{Left-child}(i) = 2i + 1$
  - $\text{Right-child}(i) = 2i + 2$ .



11

11

## Heap Operations : sift-up and sift-down

- All binary heap operations are built from the two fundamental operations *sift-up* and *sift-down*:
  - *sift-up*( $i$ ) : Repeatedly swap element  $A[i]$  with its parent as long as  $A[i]$  violates the heap property with respect to its parent (i.e., as long as  $A[i] < A[\text{parent}(i)]$ ).
  - *sift-down*( $i$ ) : As long as  $A[i]$  violates the heap property with one of its children, swap  $A[i]$  with its smallest child.
- Both operations run in  $O(\log n)$  time since the height of an  $n$ -element heap is  $O(\log n)$ .
- In some other places, *sift-down* is called *heapify*, and *sift-up* is known as *up-heap*.

12

12


## Implementing Heap Operations Using sift-up and sift-down

- The remaining operations are now easy to implement in terms of *sift-up* and *sift-down*:
  - *insert* : place new element in  $A[n+1]$ , then *sift-up*( $n+1$ ).
  - *remove-min* : swap  $A[n]$  and  $A[1]$ , then *sift-down*(1).
  - *decrease-key*( $i, \Delta k$ ) : decrease  $A[i]$  by  $\Delta k$ , then *sift-up*( $i$ ).
  - *increase-key*( $i, \Delta k$ ) : increase  $A[i]$  by  $\Delta k$ , then *sift-down*( $i$ ).
  - *delete*( $i$ ) : swap  $A[i]$  with  $A[n]$ , then *sift-up*( $i$ ), *sift-down*( $i$ ).
- All of these clearly run in  $O(\log n)$  time.
- General idea: modify the heap, then fix any violation of the heap property with one or two calls to *sift-up* or *sift-down*.

13

13

## Caveat: You Can't Easily Find Elements In Heaps (Except the Min)



Each record in the data structure keeps a pointer to the physical element of data it represents, and each element of data maintains a pointer to its corresponding record in the data structure.

14

14

## Building a Binary Heap


- We could build a binary heap in  $O(n \log n)$  time using  $n$  successive calls to *insert*.
- Another way to build a heap: start with our  $n$  elements in arbitrary order in  $A[0..n-1]$ , then call *sift-down*( $i$ ) for  $i = n-1$  down to 0.
  - Remarkable fact #1: this builds a valid heap!
  - Remarkable fact #2: this runs in only  $O(n)$  time!

15

15

## Bottom-Up Heap Construction

- The key property of *sift-down* is that it fixes an isolated violation of the heap property at the root:



- Using induction, it is now easy to prove that our “bottom-up” construction yields a valid heap.

16

16



## Bottom-Up Heap Construction

- To analyze the running time of bottom-up construction, note that:
  - At most  $n$  elements reside in the bottom level of the heap. Only 1 unit of work done to them by *sift-down*.
  - At most  $n/2$  elements reside in the 2<sup>nd</sup> lowest level, and at most 2 units of work are done to each of them.
  - At most  $n/4$  elements reside in the 3<sup>rd</sup> lowest level, and at most 3 units of work are done to them.
- So total time  $\leq T = n + 2(n/2) + 3(n/4) + 4(n/8) + \dots$   
(for simplicity, we carry the sum out to infinity, as this will certainly give us an upper bound).
- Claim:  $T = 4n = O(n)$

17

17

## “Shifting” Technique for Sums

$$\begin{array}{rcl}
 T & = & n + 2(n/2) + 3(n/4) + 4(n/8) + \dots \\
 - \quad T/2 & = & \quad n/2 + 2(n/4) + 3(n/8) + \dots \\
 \hline
 T/2 & = & n + n/2 + n/4 + n/8 + \dots
 \end{array}$$

Applying the same trick again:

$$\begin{array}{rcl}
 T & = & 2n + n + (n/2) + (n/4) + \dots \\
 - \quad T/2 & = & \quad n + (n/2) + (n/4) + \dots \\
 \hline
 T/2 & = & 2n
 \end{array}$$

18

18

## Heapsort


- Any priority queue can be used to sort. Just use  $n$  *inserts* followed by  $n$  *remove-mins*.
- The binary heap gives us a particularly nice way to sort in  $O(n \log n)$  time, known as **heapsort**:
  - Start with an array  $A[0..n-1]$  of elements to sort.
  - Build a heap (bottom up) on  $A$  in  $O(n)$  time.
  - Call *remove-min*  $n$  times.
  - Afterwards,  $A$  will end up reverse-sorted (it would be forward-sorted if we had started with a “max” heap)

19

19

## Recall: An Alternative Method With Simpler(?) Structure...

- Suppose we store our priority queue in a “heap-ordered” binary tree.
  - Heap property:  $\text{parent} \leq \text{child}$ .
  - Each node maintains a pointer to its left child and right child.
  - The tree is not necessarily “balanced”. It could conceivably be nothing more than a single sorted path.
  - No longer easily mapped to an array, as with a binary heap.



20

20


## All You Need is Merge...

- Suppose we can *merge* two heap-ordered trees in  $O(\log n)$  time.
- All priority queue operations now easy to implement in  $O(\log n)$  time!
  - *insert*: merge with a new 1-element tree.
  - *remove-min*: remove root, merge left & right subtrees.
  - *decrease-key & increase-key*: delete + re-insert
  - *delete*: replace with merge of two child subtrees

21

21

## Merging Two Heap-Ordered Trees (Null Path Merging Viewpoint)




22

22

## Merging Two Heap-Ordered Trees (Recursive Viewpoint)

- Take two heap-ordered trees  $h_1$  and  $h_2$ , where  $h_1$  has the smaller root.
- Clearly,  $h_1$ 's root must become the root of the merged tree.
- To complete the merge, recursively merge  $h_2$  into either the left or right subtree of  $h_1$ :



- As a base case, the process ends when we merge a heap  $h_1$  with an empty heap, the result being just  $h_1$ .

23

23

## Running Time Analysis


- The time required to merge two heaps along null paths is proportional to the combined lengths of these paths.
- So all we need is a method to find “short” null paths and we will have an efficient merging algorithm.
- Note that every  $n$ -node binary tree has a null path of length  $O(\log n)$ .
- There are many ways to find short null paths, each of which leads us to a different mergeable heap data structure...

24

24

## Recall: Skew Heaps

- Merge towards  $h_1$ 's "preferred child", then toggle preferred child pointer (so alternate merging into left and right subtrees).
  - Equivalently: Always merge into R, but afterwards just swap  $h_1$ 's children.
  - Equivalently (but more confusing...): Merge two heaps along their "right spines", then walk back up the right spine of the result, and for each element except the lowest, swap left and right children.
- Remarkably, this makes merge run in just  $O(\log n)$  amortized time!




25

25

## Size-Augmented Mergeable Heaps

- Let's try to remove the randomness from our preceding approach...
  - Augment each element with the size of its subtree.
  - Now we can find a null path of length  $O(\log n)$  by repeatedly stepping to whichever child has smaller size.
  - Each step reduces the size of our current subtree by at least a factor of 2.

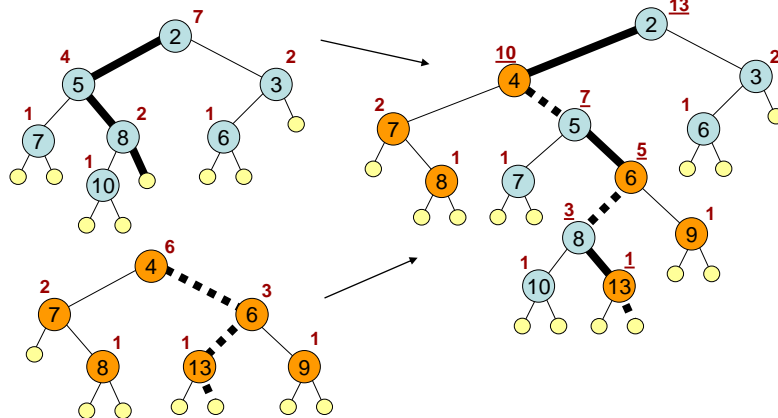


26

26

## Updating Augmented Information After Merging

- When we merge two heaps, we walk back up the merge path and update augmented information.




27

27

## Trouble with Decrease-Key Motivating Lazy Delete / Garbage Collection

- Recall: *decrease-key*(*e*) removes the subtree rooted at element *e*, decreases *e*'s key, and then merges the result back into the main tree.
- When we remove *e*'s subtree, we need to update the augmented size information along the path from the root down to *e*.
- However, if *e* is deep in the tree, this can be too expensive!
- This also affects *delete* and *increase-key*, since these are built using *decrease-key*.




28

28

## Augmenting with Null Path Lengths

- The **null path length** of element  $e$ ,  $npl(e)$ , is the shortest distance from  $e$  down to an empty space at the bottom of  $e$ 's subtree.
- Suppose we augment every element  $e$  in our heap with  $npl(e)$ .
- Since  $npl(\text{root}) = O(\log n)$ , we can find a null path of length  $O(\log n)$  by repeatedly stepping to a child with the smaller null path length.
- This allows us to *merge* in  $O(\log n)$  time.




29

29

## Leftist Heaps

- **Leftist property:**  $npl(\text{left}(e)) \geq npl(\text{right}(e))$  for all elements  $e$ .
- A **leftist heap** is a heap-ordered leftist tree. Each element in a leftist heap is augmented with its null path length.
- The shortest null path in a leftist tree (of length  $O(\log n)$ ) is its "right spine", so we can merge two leftist heaps in  $O(\log n)$  time by merging their right spines together:



30

30

## Restoring the Leftist Property


- After merging, we walk back up the right spine of the merged heap...
  - ... recalculating  $npl(e)$  for each element  $e$ .
  - ... and swapping the left and right subtrees of any element that now violates the leftist property.
- Therefore, *merge* (hence *insert* and *remove-min*) all take  $O(\log n)$  time on a leftist heap.
- The same problem exists with *decrease-key*, *delete*, and *increase-key* as with size-augmented and  $npl$ -augmented heaps though.
- **In fact, the leftist heap is really nothing more than an  $npl$ -augmented heap where we always treat the child of smaller  $npl$  as the left child.**

31

31

## The Randomized Mergeable Binary Heap

- Perhaps the simplest possible idea: choose null paths at random! (i.e., starting from root, repeatedly step left or right, each with probability  $\frac{1}{2}$ ).
- In terms of our recursive outlook for merging  $h_1$  and  $h_2$  ( $h_1$  having the smaller root) this corresponds to the following simple procedure:



- Remarkably, this trivial procedure merges any two heaps in  $O(\log n)$  time with high probability!

32

32



## Definition : “With High Probability”

- We say an algorithm with input size  $n$  runs in  $O(\log n)$  time **with high probability** if we can find a constant  $k$  such that

$$\Pr[\text{running time} > k \log n] \leq X,$$

where  $X$  is a sufficiently small number.  
but how small should  $X$  be?...

33

33

## Definition : “With High Probability”

- We say an algorithm with input size  $n$  runs in  $O(\log n)$  time **with high probability** if we can find a constant  $k$  such that

$$\Pr[\text{running time} > k \log n] \leq X,$$

where  $X$  is a sufficiently small number.  
but how small should  $X$  be?...

$$\Pr[\text{Struck by lightning in a year}] \approx 10^{-6}$$

34

34

## Definition : “With High Probability”

- We say an algorithm with input size  $n$  runs in  $O(\log n)$  time **with high probability** if we can find a constant  $k$  such that

$$\Pr[\text{running time} > k \log n] \leq X,$$

where  $X$  is a sufficiently small number.  
but how small should  $X$  be?...

$$\Pr[\text{Struck by lightning in a year}] \approx 10^{-6}$$

$$\Pr[\text{Bitten by shark in a year}] \approx 10^{-7}$$

35

35

## Definition : “With High Probability”

- We say an algorithm with input size  $n$  runs in  $O(\log n)$  time **with high probability** if we can find a constant  $k$  such that

$$\Pr[\text{running time} > k \log n] \leq X,$$

where  $X$  is a sufficiently small number.  
but how small should  $X$  be?...

$$\Pr[\text{Struck by lightning in a year}] \approx 10^{-6}$$

$$\Pr[\text{Bitten by shark in a year}] \approx 10^{-7}$$

$$\Pr[\text{Hit by a meteor in a year}] \approx 10^{-10}$$

36

36

## Definition : “With High Probability”

- We say an algorithm with input size  $n$  runs in  $O(\log n)$  time **with high probability** if we can find a constant  $k$  such that

$$\Pr[\text{running time} > k \log n] \leq X,$$

where  $X$  is a sufficiently small number.  
but how small should  $X$  be?...

$$\Pr[\text{Struck by lightning in a year}] \approx 10^{-6}$$

$$\Pr[\text{Bitten by shark in a year}] \approx 10^{-7}$$

$$\Pr[\text{Hit by a meteor in a year}] \approx 10^{-10}$$

$$\Pr[\text{All 3 in the same year!}] \approx 10^{-23}$$

37

37

## Definition : “With High Probability”

- We say an algorithm with input size  $n$  runs in  $O(\log n)$  time **with high probability** if, for any constant  $c$ , we can find another constant  $k$  such that

$$\Pr[\text{running time} > k \log n] \leq 1 / n^c.$$

- That is, the probability we fail to run in  $O(\log n)$  time is at most  $1 / n^c$ , for any constant  $c$  of our choosing (as long as we choose a sufficiently large hidden constant in the  $O(\log n)$  notation).
- We'll discuss and motivate this definition in more detail later in the course.

38

38

## Example: Boosting Success Probability via Independent Repetition

- Take a randomized algorithm that fails with probability  $\leq \frac{1}{2}$ . (and that we can detect failures).
- Run it  $k$  times: probability of failure drops to  $\leq 1/2^k$ .
- Run it  $k = c \log n$  times: the probability of failure drops to  $\leq 1/n^c$  (i.e., the algorithm succeeds with high probability!)

39